


Introduction to Apache Kafka



Thessaloniki not-only-Java Meetup

Dimitris Kontokostas @Diffbot

About me...

Data geek,

Software engineer

Open source enthusiast,

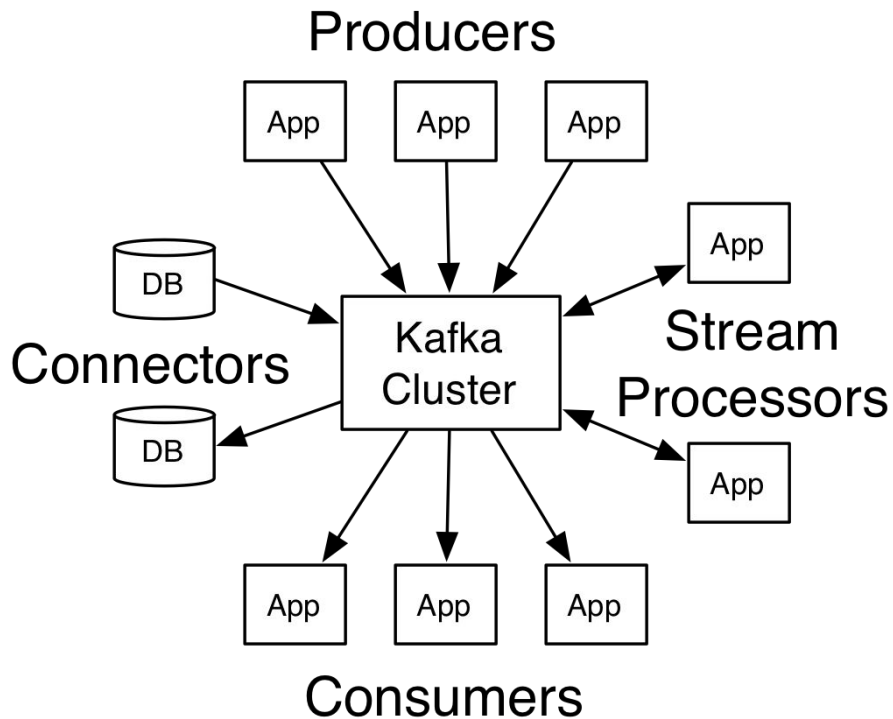
Remote Worker

Currently working @ [Diffbot](#)

Summary

- The Log
- Kafka 101
- KStream
- KSQL
- Connectors
- Schemas

Let's keep this interactive...



Kafka is

A distributed streaming platform

A distributed publish-subscribe messaging system/queue

A distributed, immutable event logging data store ^(*)

The Log


What every software engineer should know about real-time
data's unifying abstraction

A great post by Jay Kreps (creator of Kafka)

Parts of the post are used in this presentation


What is a log

- Records **events**
- **What** happened & **when**
- **Append only**, left-to-right
- Timestamps



Logs in databases

- As internal structures to facilitate ACID
- Gradually getting more exposed
- As a data subscription mechanism





Logs in distributed systems

Ordering

Replication

Active-active vs Active-passive

- E.g. "+5", "-2", "*4"
- Vs "0", "5", "3", "12"



Events & Tables

Two sides of the same coin



events & tables

1. **User_add** {id:1, email: "joh@doe.com"}
2. **User_change** {id:1 email: "john@doe.com"}
3. **User_add** {id:2, email: "mary@doe.com"}
4. **User_change** {id:2, email: "mary@doe.gr"}
5. **User_add** {id:3, email: "peter@doe.com"}

Table User

id	email
1	?
2	?
3	?

events & tables

1. **User_add** {id:1, email: "[joh@doe.com](mailto:john@doe.com)"}
2. **User_change** {id:1 email: "john@doe.com"}
3. **User_add** {id:2, email: "mary@doe.com"}
4. **User_change** {id:2, email: "mary@doe.gr"}
5. **User_add** {id:3, email: "peter@doe.com"}

Table User

id	email
1	john@doe.com
2	mary@doe.gr
3	peter@doe.com


events & tables

1. **User_add** {id:1, email: "joh@doe.com"}
2. **User_change** {id:1 email: "john@doe.com"}
3. **User_add** {id:2, email: "mary@doe.com"}
4. **User_change** {id:2, email: "mary@doe.gr"}
5. **User_add** {id:3, email: "peter@doe.com"}
6. **User_remove**{id:1 }

Table User

id	email
1	john@doe.com
2	mary@doe.gr
3	peter@doe.com

events & tables



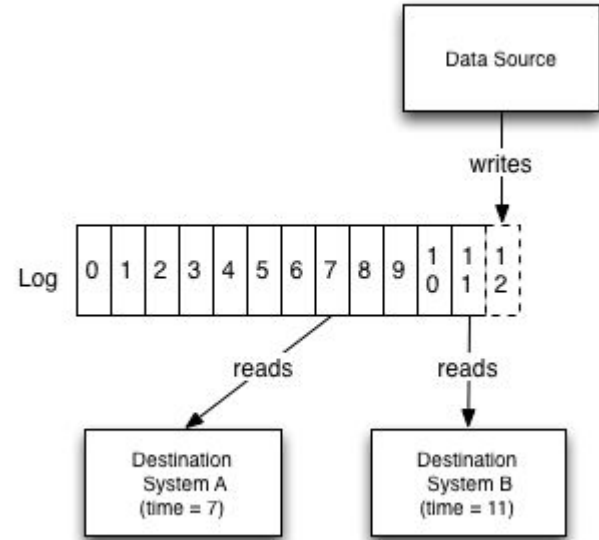


I have a cool idea!!!

Report users with for more than 2
email edits per day as suspicious

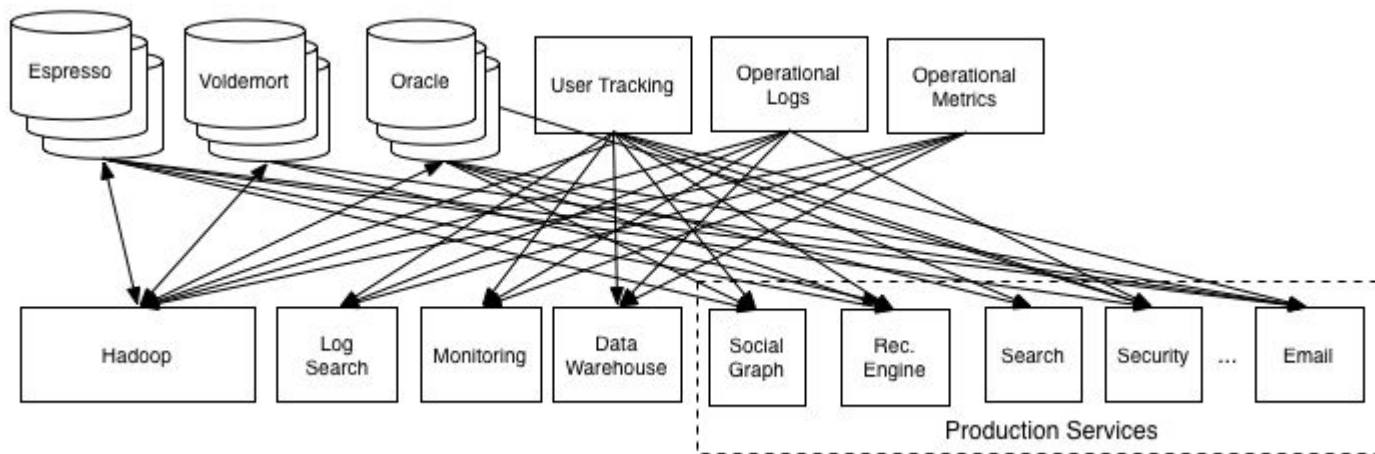
What if we also did X...

Decoupling from the table

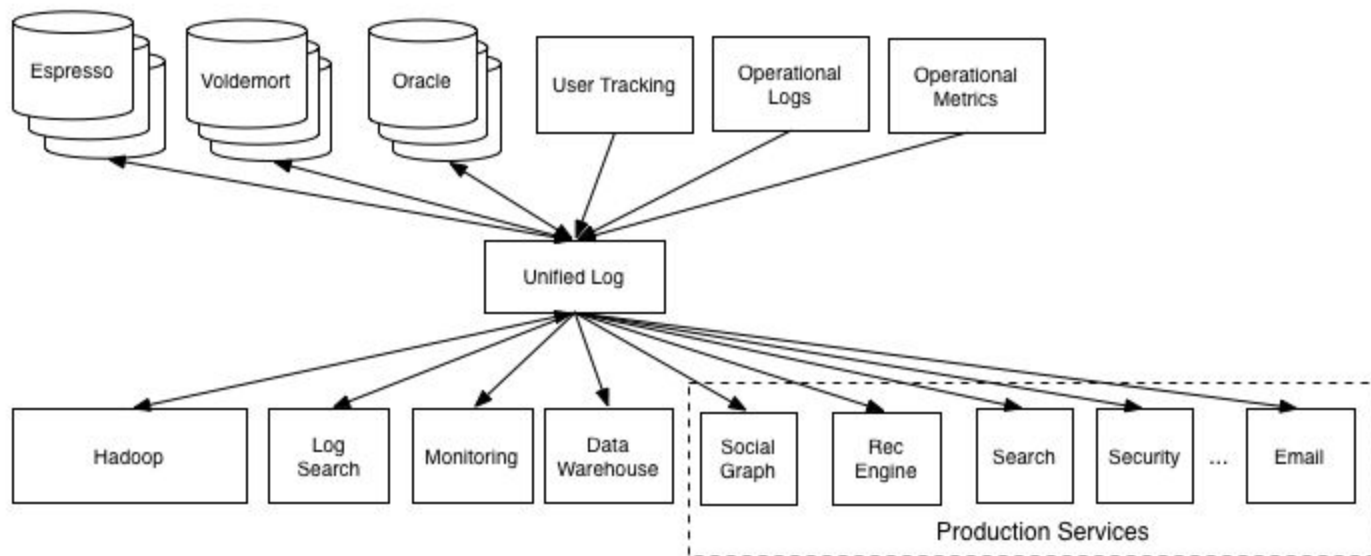


Tables are opinionated views of your data
The later you can defer your opinion, the better...

LinkedIn before Kafka



LinkedIn after Kafka



Kafka Basics

Basic terminology

Kafka maintains feeds of **messages** in categories called **topics**

Kafka Producers are processes that **publish messages** on a kafka topic

Kafka Consumers are processes that **subscribe to topics** and process the feed of published messages


Messages can be in any form, text, json, binary etc.

Kafka brokers are the servers that comprise the **kafka cluster**

Topics & Logs

1 Topic \approx 1 Partitioned Log
(with configurable retention period)

Anatomy of a Topic



What is the difference with “The Log”?

Ordering...

Kafka Messages

Anything can be a message

=> Strings, JSON, XML, binary formats



Messages are preferably small (< few Kb)

=> Big messages affect brokers & throughput

=> Reference big messages as external resources

Message IDs define which partition the message is stored


=> Default behavior / can be manually overwritten



Producers & Consumers

Message Delivery Semantics

- => At least once
- => At most once
- => Exactly once



Multiple (independent) Consumers can read from the same topic

- => Consumers manage their own offset (stored on Kafka)
- => Messages remain on Kafka

Consumer Scaling

Every Consumer has a **Consumer ID**

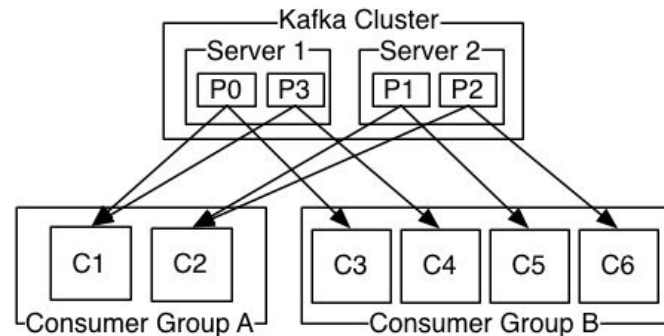

=> Kafka keeps track of the offsets

=> rejoining will continue from latest offset

Consumers with the same ID
belong to the same **Consumer Group**

Consumers from the same group
read messages from **different partitions**

=> Topic partition size is the hard limit



KStreams/KSQL

Print topic messages (Consumer API)

```
public class KafkaConsumerExample {  
    ...  
    static void runConsumer() throws InterruptedException {  
        final Consumer<Long, String> consumer = createConsumer();  
        while (true) {  
  
            final ConsumerRecords<Long, String> consumerRecords = consumer.poll();  
  
            consumerRecords.forEach(record -> {  
                System.out.printf( "Consumer Record: (%d, %s, %d, %d)\n" ,  
                    record.key(), record.value(), record.partition(), record.offset());  
            });  
  
            consumer.commitAsync();  
        }  
        consumer.close();  
    }  
}
```

Print topic messages (KStream API)

```
static void createWordCountStream(final StreamsBuilder builder) {  
  
    final KStream<String, String> textLines = builder.stream(inputTopic);  
  
    textLines  
        .forEach(value -> System.out.println(value));  
}
```

KStream API

Abstraction on top of classic Kafka API

Provide a Java Stream-like API on Kafka

=> has a source, sink & stream processors

=> aggregations, windowing, KTables / state

Source processor


=> consumes on or more topics, forwards downstream

Stream processors

=> apply transformations, aggregations, etc

Sink Processors

=> final processors, stores results on a topic or KTable



KTables

Read-only & fault-tolerant state store

Can be in memory (on consumer) or global (persistent on the cluster)

RocksDB is a default implementation

Example: keep an up-to-date table of `[user id, email]`

Aggregations & Windowing

Example: Users with more than 2 email changes in 24 hours

```
final KTable<Windowed<String>, Long> anomalousUsers = views
    .groupByKey()
    .windowedBy(TimeWindows.of(Duration.ofDays(1)))
    .count()
    .filter((windowedUserId, count) -> count >= 2);

final KStream<String, Long> anomalousUsersForConsole = anomalousUsers
    .toStream()
    .filter((windowedUserId, count) -> count != null)
    .map((windowedUserId, count) -> new KeyValue<>(windowedUserId.toString(), count));

anomalousUsersForConsole.to("AnomalousUsers", Produced.with(stringSerde, longSerde));
```

KSQL

```
CREATE TABLE users_view AS  
  SELECT * FROM USERS;
```

```
CREATE STREAM suspicious_users AS  
  SELECT id, count(*)  
  FROM USERS  
  WINDOW SESSION (1 DAY)  
  GROUP BY id;  
  HAVING count(*) >= 2;
```

Disclaimer: most probably has typos !!!

Connectors & Schema Registries

Big library of source & sink Kafka connectors

- => S3, most DBs, ES, HDFS, other streaming frameworks
- => e.g. write this topic to MySQL table X
- => Kafka Connect API for writing a custom one

Schemas always change, not only in RDBMS :)

- => Message body freedom has a cost
- => Schema registries come to the... hmm... help...
- => data validation, de/serialization, backwards compatibility

Summary

Logs are everywhere

Logs help us defer data schema decision

Kafka is an event logging data store

Has a great ecosystem of libraries & integrations

Thank you for your attention

Questions?